

Reading: Chapter 07 of [Chiusano and Bjarnason 2014]

Purely Functional Parallelism

Lecture 060 of Advanced
Programming

March 7, 2018

Andrzej Wasowski & Zhoulai Fu

IT University of Copenhagen

How today's class differs from your PCPP (Practical Concurrent and Parallel Programming) class

- ❖ PCPP is about writing efficient concurrent and parallel software, using existing API in Java
- ❖ Today's class is about *designing* such API, under the functional paradigm
- ❖ A minor note about concurrency vs parallelism:
 - ❖ Concurrency describes a *problem* — that things need to happen together
 - ❖ Parallelism describes a *solution* — *that is* based on multiple threads / CPUs
 - ❖ Other solutions exist; we focus on parallelism today

Prerequisite 1: ExecutorService and Future API

```
class ExecutorService {  
    def submit[A] (a: Callable[A]): Future[A]  
}  
trait Future[A] {  
    def get: A  
}
```

```
public interface ExecutorService  
extends Executor
```

An **Executor** that provides methods to manage termination and methods that can produce a **Future** for tracking progress of one or more asynchronous tasks.

DEMO (if time allows)

```
val es=Executors.newWorkStealingPool()  
val fut=es.submit(new Callable[Int] {  
    override def call(): Int = (1 to 10).sum  
})  
val res=fut.get
```


Prerequisite 2: Strict and Lazy Functions

- ❖ Strict Evaluation (by-value): function argument evaluated before entering the function
- ❖ Lazy evaluation, (by-name, by-need)

DEMO (if time allows)

```
def f_lazy(i: => Int) = {println("Entering call");i}
def f_strict(i: Int) = {println("Entering call");i}

f_lazy{println("Calling"); 42} // Entering call, calling
f_strict{println("Calling"); 42} // Calling, Entering call
```


Designing the API

- No right answers in design
- You will see a collection of design choices
- You are to understand their trade-offs, and think critically.

Why not use Java Thread

```
trait Runnable { def run: Unit }
```

```
class Thread(r: Runnable) {  
  def start: Unit  
  def join: Unit  
}
```

**Begins running x in
a separate thread.**

**Blocks the calling thread
until x finishes running.**

- ❖ **Side-effects are evil when it comes to program reasoning:** if we want to get any information out of a `Runnable`, it has to have some side effect, like mutating some state that we can inspect. This is bad for compositionality—we can't manipulate `Runnable` objects generically since we always need to know something about their internal behavior.
- ❖ **Thread is low-level and low-efficient:** `Thread` maps directly onto operating system threads, which are a scarce resource. It would be preferable to create as many “logical threads” as is natural for our problem, and later deal with mapping these onto actual OS threads.

Design Goals

- ❖ **Pure:** function return the same value for the same input, without observable side effects
- ❖ **High-level:** having the capability to write something like foldleft, as in sequential programs.

```
def sum(ints: Seq[Int]): Int =  
    ints.foldLeft(0) ((a,b) => a + b)
```

Design Methodologies

- ❖ Start from simple examples
- ❖ Try - Challenge - Refine

Example: Summing a list with divide-and-conquer

Divides the sequence in half using the `splitAt` function.

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints.headOption getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    sum(l) + sum(r)  
  }
```

IndexedSeq is a superclass of random-access sequences like `Vector` in the standard library. Unlike lists, these sequences provide an efficient `splitAt` method for dividing them into two parts at a particular index.

`headOption` is a method defined on all collections in Scala. We saw this function in chapter 4.

Recursively sums both halves and adds the results together.

❖ Listing 7.1, [Chiusano et al]

The Making of a Parallel Sum (1 - try)

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints headOption getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    val sumL: Par[Int] = Par.unit(sum(l))  
    val sumR: Par[Int] = Par.unit(sum(r))  
    Par.get(sumL) + Par.get(sumR)  
  }
```

Computes the left half in parallel.

Computes the right half in parallel.

Extracts both results and sums them.

- ❖ Need a data type to contain parallel computation results: **Par[A]**
- ❖ Need a function to evaluate a computation in a separate thread
 - ❖ **Par.unit** (a: =>A): Par[A]
- ❖ Need another function to extract a result from a Par[A]:
 - ❖ **Par.get** [A] (a: Par[A]):A

The Making of a Parallel Sum (1 - problem)

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints.headOption.getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    val sumL: Par[Int] = Par.unit(sum(l))  
    val sumR: Par[Int] = Par.unit(sum(r))  
    Par.get(sumL) + Par.get(sumR)  
  }
```

Computes the left half in parallel.

Computes the right half in parallel.

Extracts both results and sums them.

- ❖ For the sake of parallelization, Par.unit has to delay the computation until Par.get
- ❖ Problem: The whole computation is still sequential because “+” is strict

The Making of a Parallel Sum: (2 - try)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.size <= 1)  
    Par.unit(ints.headOption.getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(sum(l), sum(r))(_ + _)  
  }
```

- ❖ Par.map2 is a new higher-order function for combining the result of two parallel computations.
- ❖ Q: What is its signature?
- ❖ A: `Par.map2[A,B,C] (a: Par[A], b: Par[B]) (f: (A,B) => C): Par[C]`
- ❖ Q: Should Par.map2 be lazy or strict?
- ❖ A: :If it is strict, we'll strictly construct the entire left half of the tree of summations first before moving on to (strictly) constructing the right half ==> Let Par.map2 be lazy

The Making of a Parallel Sum: (2 - problem)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.size <= 1)  
    Par.unit(ints.headOption getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(sum(l), sum(r))(_ + _)  
  }
```

- ❖ Q: Do we always want to evaluate the two arguments to `Par.map2` in parallel?
- ❖ A: : Probably not. Consider `Par.map2(Par.unit(1), Par.unit(2))(_+_)`. The overhead for thread creation/management is swamping any tiny gains from parallelization.
- ❖ Problem: This API is ver inexplicit about when computations gets forked off the main thread — the programmer cannot specify where this forking should occur.

The Making of a Parallel Sum: (3 - try)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.length <= 1)  
    Par.unit(ints.headOption getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(Par.fork(sum(l)), Par.fork(sum(r))) (_ + _)  
  }
```

- ❖ `Par.fork[A](a: => Par[A]): Par[Int]` runs *a* in a separate logical thread
- ❖ With *Par.fork*, you can make *Par.map2* strict, leaving it up to the programmer to wrap arguments if they want

The Making of a Parallel Sum: (3 - problem)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.length <= 1)  
    Par.unit(ints.headOption getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(Par.fork(sum(l)), Par.fork(sum(r))) (_ + _)  
  }
```

- ❖ Problem: If *Par.fork* runs *sum(l)* or *sum(r)* immediately in a separate thread, the thread pool (or whatever resource we use to implement the parallelism) must be globally accessible and properly initialized wherever we want to call — meaning that we lose the ability to control the parallelism strategy used for different parts of our program.
- ❖ Although there's nothing inherently wrong with having a global resource for executing parallel tasks, it seems more appropriate to give *get* the responsibility of creating threads and submitting execution tasks.

The Making of a Parallel Sum: (final)

- ❖ We let *fork* hold on to its unevaluated argument until later. It takes an unevaluated $Par[A]$ and marks it for concurrent evaluation later
- ❖ In this model, $Par[A]$ holds a *description* of a parallel computation that gets *interpreted* at a later time by something like the *get* function
- ❖ How to implement it?

Implementation: reused API from Java

```
class ExecutorService {  
  def submit[A](a: Callable[A]): Future[A]  
}  
trait Callable[A] { def call: A }  
trait Future[A] {  
  def get: A  
  def get(timeout: Long, unit: TimeUnit): A  
  def cancel(evenIfRunning: Boolean): Boolean  
  def isDone: Boolean  
  def isCancelled: Boolean  
}
```

- ❖ *ExecutorService* lets us submit a *Callable* value and get back a corresponding *Future*, which is a handle to a computation that's potentially running in a separate thread.
- ❖ When *Future* obtain a value from *get*, it blocks the current thread until the value is available.
- ❖ *Future* has some extra features for cancellation (throwing an exception after blocking for a certain amount of time, and so on).

Implementation: Other API

- ❖ Type alias: *type Par[A] = ExecutorService => Future[A]*
- ❖ Object Par that holds three primitive operations: *unit*, *map2*, and *fork*



Quiz

- ❖ Define $\text{map}[A,B](pa: \text{Par}[A])(f: A \Rightarrow B): \text{Par}[B]$ in terms of:
 - ❖ $\text{map2}[A,B,C] (a: \text{Par}[A], b: \text{Par}[B]) (f: (A,B) \Rightarrow C): \text{Par}[C]$

Quiz

- ❖ Define `map[A,B](pa: Par[A])(f: A => B): Par[B]` in terms of:
 - ❖ `map2[A,B,C] (a: Par[A], b: Par[B]) (f: (A,B) => C): Par[C]`
- ❖ Solution:

```
def map[A,B](pa: Par[A])(f: A => B): Par[B]  
= map2(pa, unit(()))((a,_) => f(a))
```
- ❖ The fact that we can implement `map` in terms of `map2` but not the other way around, just shows that `map2` is strictly more powerful than `map`.
- ❖ This sort of thing happens a lot when we're designing libraries—often, a function that seems to be primitive will turn out to be expressible using some more powerful primitive.

Answer

❖ Define

❖ $\text{map}[A,B](\text{pa}: \text{Par}[A])(f: A \Rightarrow B): \text{Par}[B]$

❖ in terms of:

❖ $\text{map2}[A,B,C] (a: \text{Par}[A], b: \text{Par}[B]) (f: (A,B) \Rightarrow C): \text{Par}[C]$

```
def map[A,B] (pa: Par[A]) (f: A => B): Par[B] =  
  map2(pa, unit(())) ((a,_) => f(a))
```

Laws and Properties

- ❖ Consider a unit test (incidentally of the unit function):
 - ❖ `map(unit(1)) (_ + 1) == unit(2)`
- ❖ What does equality mean on the `Par[Int]` values?
- ❖ For instance with the following definition of equality:
 - ❖ `def equal[A] (e: ExecutorService) (p: Par[A], p2: Par[A]) :Boolean = p(e).get == p2(e).get`
- ❖ But how would we test `map`? A more general test would be nice:
 - ❖ `map (pa) (f).get == f(pa.get)` for all `pa, f`
- ❖ This is no longer a unit test, but a property.
- ❖ Next week you will look into how to turn such properties into tests systematically

Takeaway

- ❖ Not only how to write a library for purely functional parallelism, but *how to approach the problem of designing a purely functional library.*